

Introduction to Threads

Mehrshad Saadatinia

Contents

- Parallelism vs. Concurrency vs. Asynchronous
- What are Threads?
- Threads vs. Processes
- Threads in different programming languages
- Threads in Java
- Basic Threads methods
- Challenges of multi-threading



Parallelism vs.
 Concurrency vs.
 Asynchronous

Can you run two tasks in parallel?

Is concurrent execution the same as parallel?

What is asynchronous then?



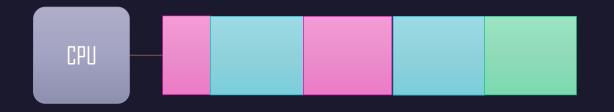
Sequential execution

One task starts after the other finishes, like programs we wrote **so far**.



Parallelism vs. Concurrency

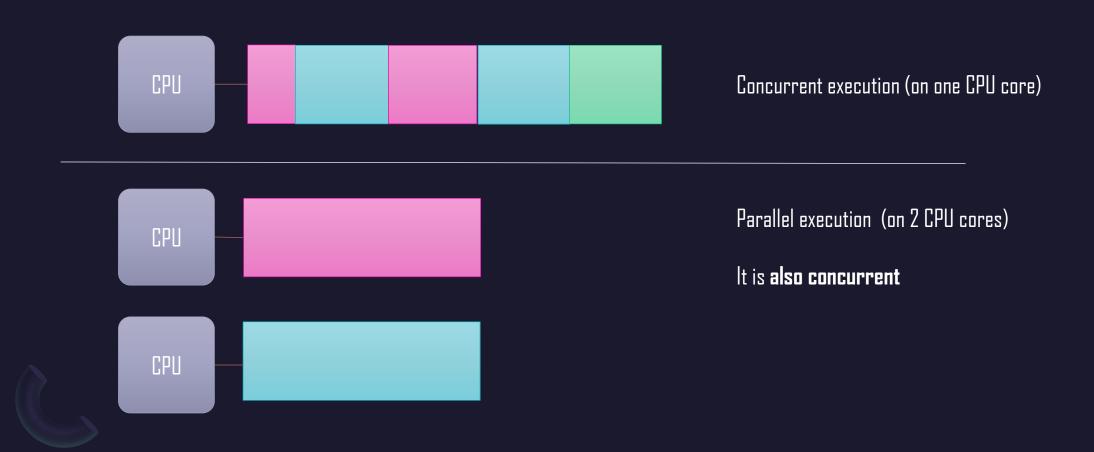
Concurrency means **executing** multiple tasks at the same time but not necessarily simultaneously.



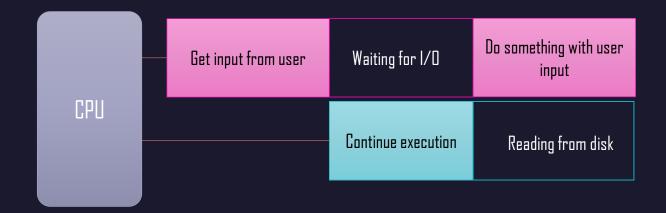
Concurrent execution (on one CPU core)

Parallelism vs. Concurrency

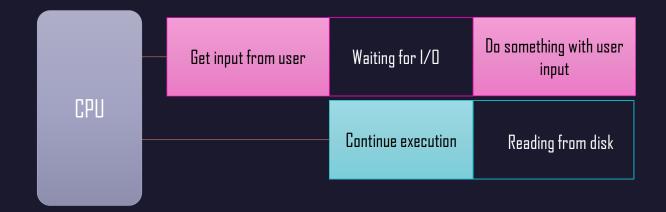
Concurrency means executing multiple tasks at the same time but not necessarily simultaneously.



Asynchronous execution



Asynchronous execution



In asynchronous execution CPU is never Idle , therefore when one task is forced to wait for something, CPU is switched to other tasks until the wait is over.

Asynchronous is also a special case of Concurrent execution

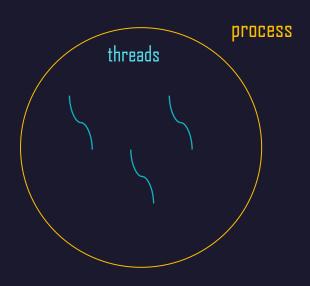


Process vs. Threads



• Any program we execute is a process

Process vs. Threads



- Any program we execute is a process
- Each process may consist of one or multiple **threads** of execution
- Each thread is created to do a task in the program
- Threads are run concurrently so multiple tasks are executed (progress) at the same time
- In some programming languages threads could be executed in parallel (like Java)

multithreading is a general programming concept and multiple programming languages implement it in different ways.





multithreading is a general programming concept and multiple programming languages implement it in different ways.



Java: Threads are run concurrently and could also run in parallel on different CPU cores



multithreading is a general programming concept and multiple programming languages implement it in different ways.



Java: Threads are run concurrently and could also run in parallel on different CPU cores



JavaScript: does not support multi-threading at all, but supports asynchronous execution



multithreading is a general programming concept and multiple programming languages implement it in different ways.



Java: Threads are run concurrently and could also run in parallel on different CPU cores



JavaScript: does not support multi-threading at all, but supports asynchronous execution



Python: threads in python are run concurrently but not in parallel



multithreading is a general programming concept and multiple programming languages implement it in different ways.



Java : Threads are run concurrently and could also run in parallel on different CPU cores



JavaScript: does not support multi-threading at all, but supports asynchronous execution



Python: threads in python are run concurrently but not in parallel, also supports asynchronous execution



Go: supports multithreading and parallelism, but it's different with Java's implementation of threads

Multi-threading in Java

Let's see how we can write programs that utilize **multiple threads**, in Java



Creating a Thread

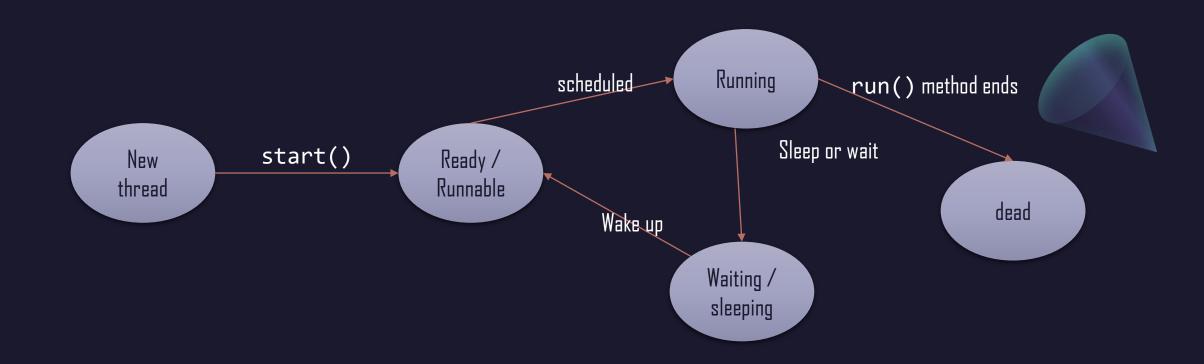
We have two choices:

- Extending Threads class (java.lang.Thread)
- 2. Implementing Runnable interface (java.lang.Runnable)

Then we implement run() method to describe our thread's tasks.

Let's see some code

Thread states



Thread methods

- Thread.currentThread(): returns the currently running thread
- Thread.sleep(millis): causes the currently executing thread to sleep for the specified number of milliseconds
- Thread.yield(): transfer control to another thread determined by scheduler
- thread1.join(): waits for thread1 to finish execution
- thread1.setPriority(Thread.MAX_PRIORITY)

Thread methods

- Thread.currentThread(): returns the currently running thread
- Thread.sleep(millis) : causes the currently executing thread to sleep for the specified number of milliseconds
- Thread.yield(): transfer control to another thread determined by scheduler
- thread1.join(): waits for thread1 to finish execution
- thread1.setPriority(Thread.MAX_PRIORITY)

```
MAX_PRIORITY = 10
MIN_PRIORITY = 1
NORM_PRIORITY= 5
```

Thread methods (2)

• thread1.setDaemon(true) : we can make thread a daemon thread (running in the background)

Example for daemon thread : garbage collector



Like anything in computer world, multi-threading has it's own challenges



critical sections

Sections in a program which we don't want multiple threads to enter at the same time

- If one threads entered critical section, no other threads must be allowed to enter
- Second thread must wait until the first one exits critical section

Why?

critical sections

Sections in a program which we don't want multiple threads to enter at the same time

- If one thread entered critical section, no other threads must be allowed to enter
- Second thread must wait until the first one exits critical section

Why?

- Shared resources:
 an entity (object, variable, ...) used in multiple threads in a shared manner
- Race condition: multiple threads trying to access a shared resource and at least one of them trying to change it

```
//shared by both threads
void critical(){
int a = 5;
if (a == 6) {
   print("this must be printed");
else if (a < 7){
   a++;
 else {
  throw error;
```

```
//shared by both threads
void critical(){
int a = 5;
 if (a == 6) {
   print("this must be printed");
 else if (a < 7){
   a++;
 else {
  throw error;
thread1 runs critical()
thread2 runs critical()
```

```
//shared by both threads
void critical(){
 int a = 5;
 if (a == 6) {
    print("this must be printed");
 else if (a < 7){
    a++;
 else {
  throw error;
thread1 runs critical()
thread2 runs critical()
```

Imagine the following scenario:

- thread I runs critical, finds a to be 5 and executes
 else if (a < 7)
- 2. Before it can increment a scheduler gives control to thread2
- 3. thread2 finds a to be still 5 and executes else if (a < 7) a++;
- 4. Now thread continues and runs a++

```
//shared by both threads
void critical(){
 int a = 5;
 if (a == 6) {
    print("this must be printed");
 else if (a < 7){
    a++;
 else {
   throw error;
thread1 runs critical()
thread2 runs critical()
```

Imagine the following scenario:

- thread I runs critical, finds a to be 5 and executes else if (a < 7)
- 2. Before it can increment a scheduler gives control to thread2
- 3. thread2 finds a to be still 5 and executes else if (a < 7) a++;
- 4. Now thread continues and runs a++

Now a is 7 and we never printed anything and If we execute critical() one more time we get an error

Solutions for critical sections

Let's look at some of the solutions proposed to solve critical section & race conditions



Synchronized blocks

```
public class BankAccount {
    private float balance;
    public synchronized void deposit(float amount) {
        balance += amount;
    }
    public synchronized void withdraw(float amount) {
        balance -= amount;
    }
}
```

Synchronized blocks

```
public class BankAccount {
    private float balance;
    public synchronized void deposit(float amount) {
        balance += amount;
    }
    public synchronized void withdraw(float amount) {
        balance -= amount;
    }
}
```

Deposit and withdraw methods are critical sections here
Two threads can't simultaneously call synchronized methods (enter synchronized block) on the same object

Inter-thread Communication

Two threads can communicate using notify() and wait() methods: Sometimes a thread needs to wait until another thread notify it to continue

There can be multiple threads on in waiting state on each object, notifyAll() will notify all of them to continue running

- *** wait and notify must be called inside a synchronized block , otherwise IllegalMonitorStateException is thrown***
- Notify and wait are native methods which means they have a low level implementation

notify() & wait()

```
synchronized( lockObject )
{
    while(! condition)
    {
       lockObject.wait();
    }

    //take the action here;
}
```

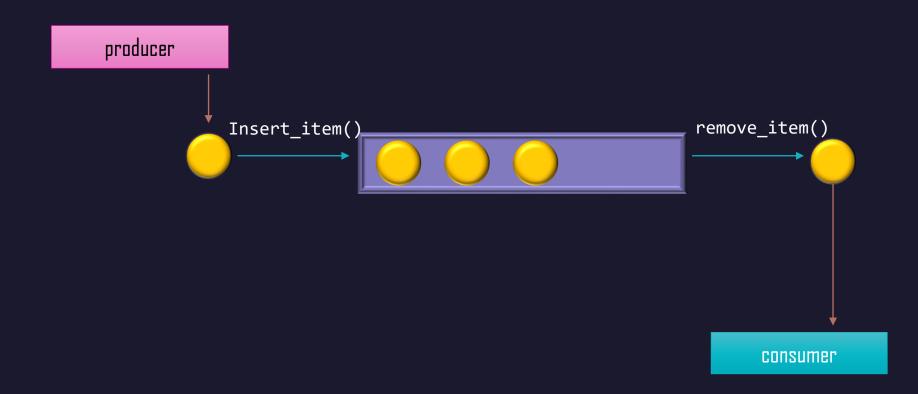
```
synchronized(lockObject)
{
    //establish_the_condition;
    lockObject.notify();
    //any additional code if needed
}
```

interrupt()

A thread might be in a waiting status, if we call interrupt on it , thread **exits waiting state** and **InterruptedException** will be thrown

If the thread is not in the sleeping or waiting state then calling the interrupt() method performs a normal behavior and doesn't interrupt the thread but sets the interrupt flag to true.

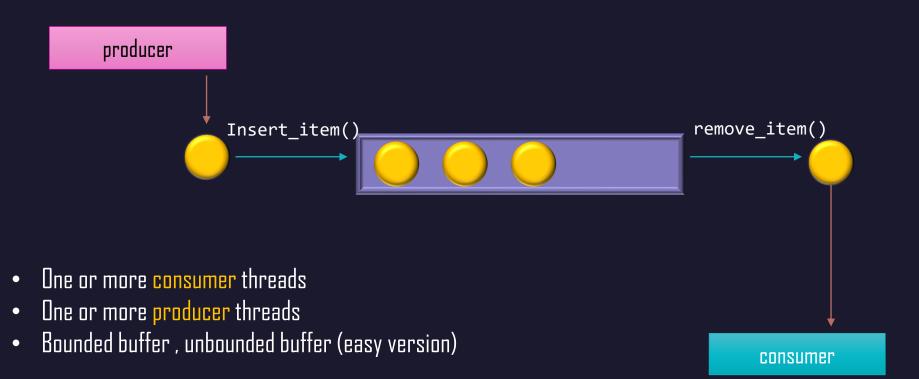
Producer Consumer problem



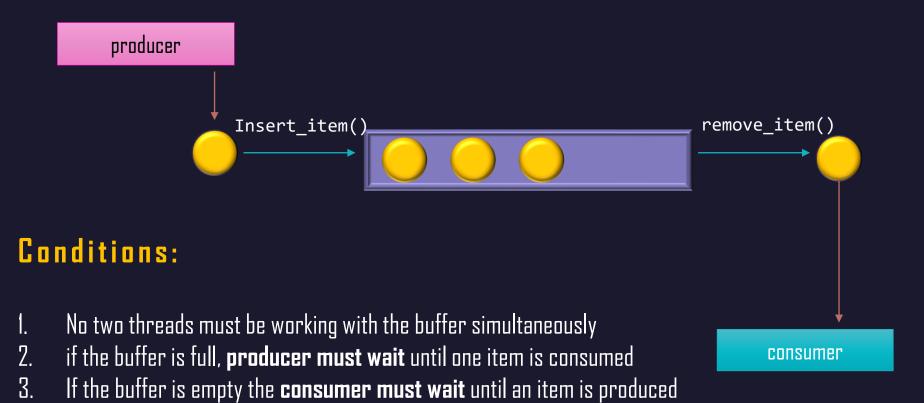
Producer Consumer problem

A traditional and important problem in concurrency settings

(Operating Systems,)



Producer Consumer problem



Lets see some code;

Thread safe Collections

When using thread safe Objects we don't to use **synchronized** because they are **thread safe**

thread safe	non thread safe
Vector	Arraylist
StringBuffer	StringBuilder
ConcurrentHashMap	HashMap

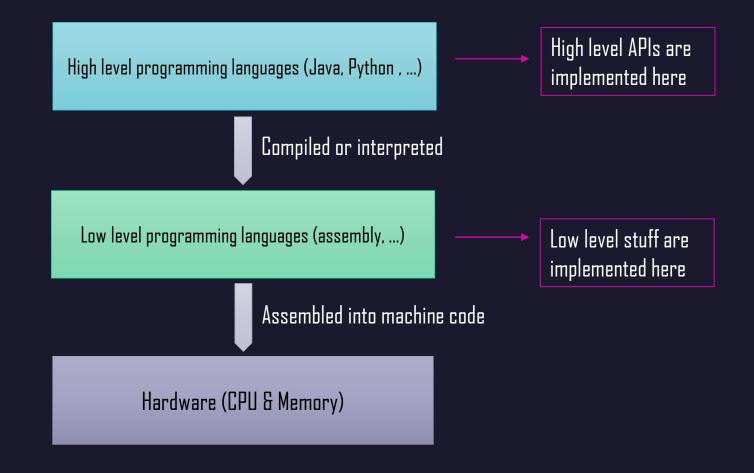
And all **immutable** Objects such as String , Integer, ...

• High-level concurrency APIs

Let's take a look at some software solutions devised for solving thread challenges and achieving **mutual exclusion***



What do we mean by high-level?



Synchronizers

Used in order to synchronize multiple thread (java.util.concurrent)

- 1. Semaphore
- 2. CountDownLatch
- 3. Exchanger
- 4. CyclicBarrier

Semaphore is a common solution to achieve mutual exclusion, and is not only limited to Java. Semaphores are widely used in Unix based **Operation Systems**

Semaphore;

Methods: acquire(), release()

A semaphore controls access to a shared resource through the use of a counter. If the counter is greater than zero, then access is allowed.

Semaphore;

Methods: acquire(), release()

A semaphore controls access to a shared resource through the use of a counter. If the counter is greater than zero, then access is allowed.

In producer /consumer problem:

```
Semaphore semaphore = new Semaphore(∅);
```

```
semaphore.acquire();
synchronized (list)
 { obj = list.remove(∅); }
```

Permits 0 threads to acquire lock

```
synchronized(list)
{ list.add(obj); }
semaphore.release();
```

Atomic classes

```
Remember this code?
void critical(){
int a = 5;
if (a == 6) {
    print("this must be printed");
else if (a < 7){
    a++;
 else {
   throw error;
```

What if we could never take execution away from Thread I while executing else if?

Atomic classes

An **atomic action** is an action which is executed as a whole, and we can not take the control away while executing it.

Java supports atomic operations on variables through atomic classes (java.util.concurrent.atomic)

Some atomic classes:

- AtomicBoolean
- AtomicInteger(extends Number)
- AtomicIntegerArray
- AtomicLongArray
- AtomicReference

And so on ...

Atomic variables are thread-safe, therefor no synchronization is needed

Lock interface

Gives us more control over how we want to execute critical sections

Lock : controlled by programmer explicitly

Synchronized: automatically locks critical section

Lock interface

Gives us more control over how we want to execute critical sections

Lock : controlled by programmer explicitly

Synchronized: automatically locks critical section

```
Lock 1 = new ReentrantLock(); // ReentrentLock implements Lock
    1.lock();
    try { ... // critical section }
finally { l.unlock(); }
```

Executor framework

Contemporary way of creating threads program determines that a part needs a new thread

In large scale enterprise application thread management is not done by programmer the programmer only defines tasks.

Executor framework is used in this case

```
interface Executor {
  void execute(Runnable command);
}
```

The executor might define a new thread for the given task , or use an existing idle thread, it depends on the type of executor and the state of **thread pool**

Executor framework

Contemporary way of creating threads program determines that a part needs a new thread

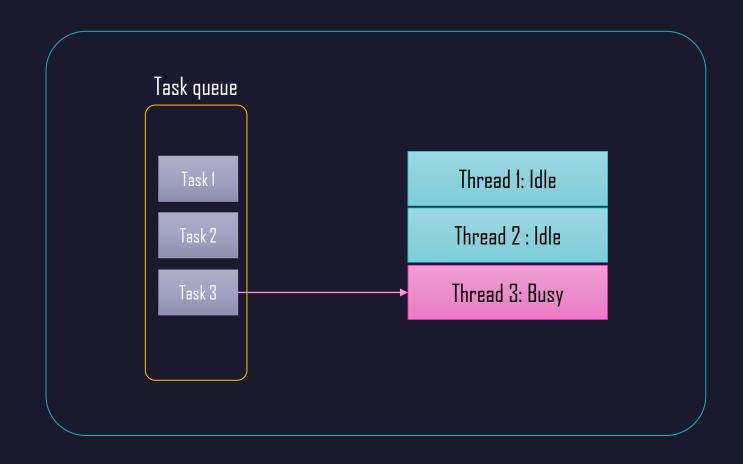
In large scale enterprise application thread management is not done by programmer the programmer only defines tasks.

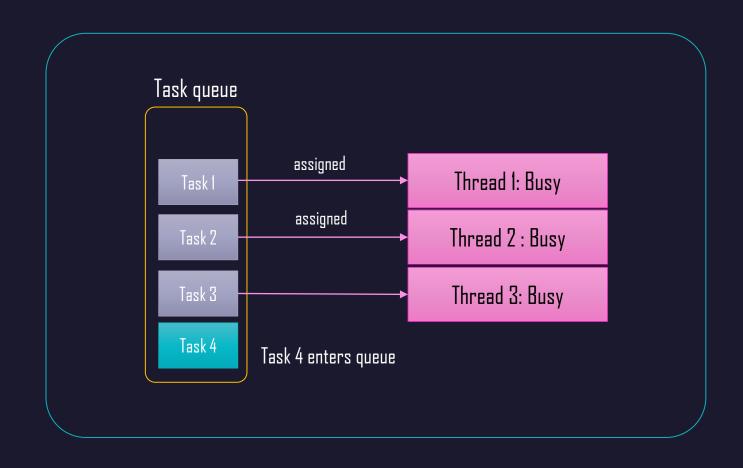
Executor framework is used in this case

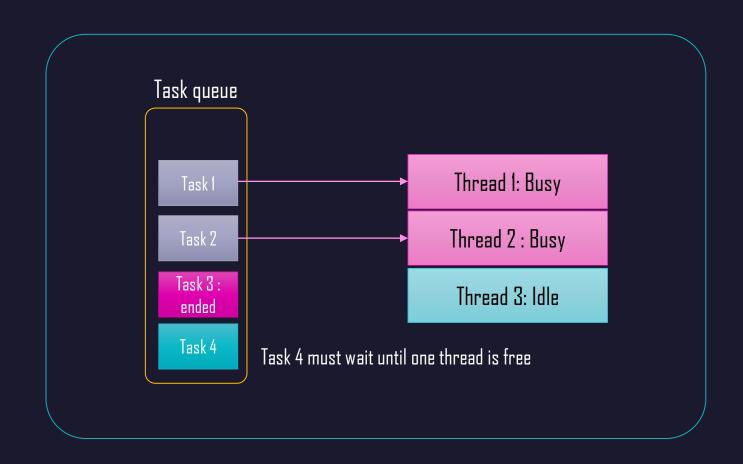
```
interface Executor {
  void execute(Runnable command);
}
```

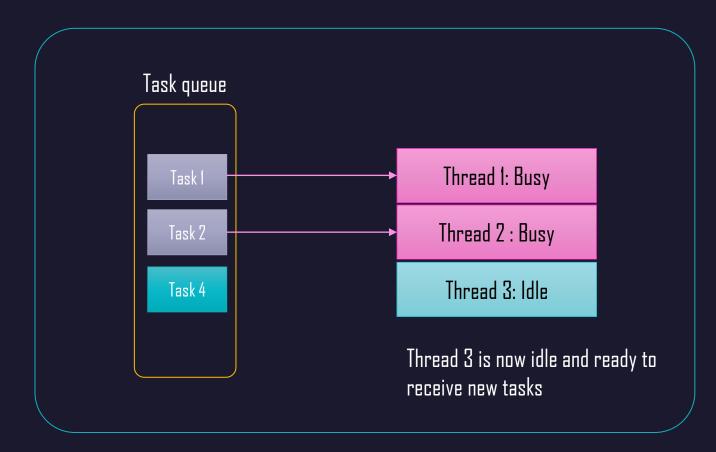
The executor might define a new thread for the given task, or use an existing idle thread, it depends on the type of executor and the state of **thread pool**

Thread what?!









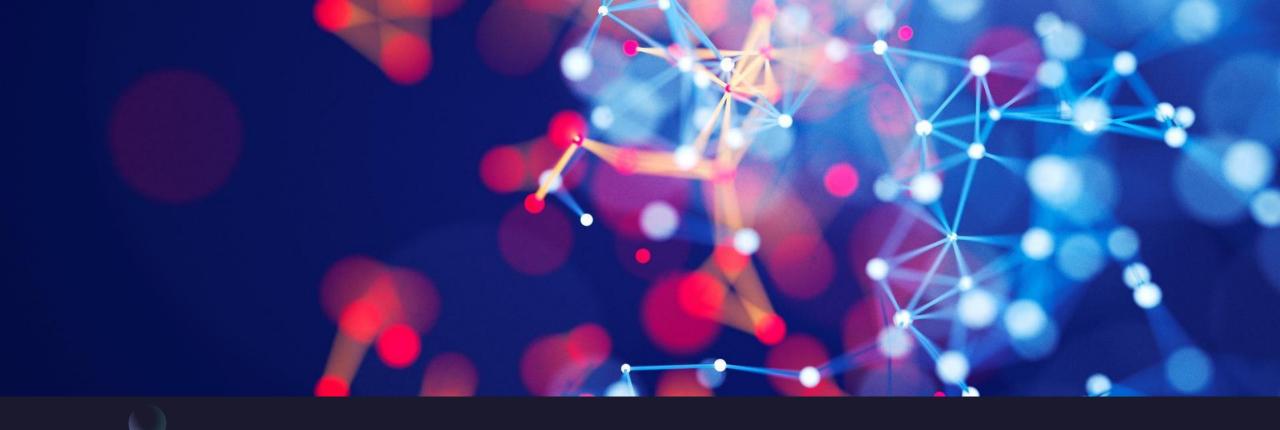
By using thread pool, executors limit the number of threads created, because creating a new thread for every single task is costly

Let's see some code;

Summary in one slide

Different types of execution concurrency Threads Threads in Java synchronized Multi-threaded programming Inter-thread Challenges Solutions communication (critical section ...) Thread safety

High level APIs Semaphore Lock interface Atomic classes Executor Thread pool framework



The end